

# Deep Learning

# Deep learning

Representation learning with a hierarchy of concepts

Those concepts are represented by layers in a neural network model

## Unsupervised models

- Autoencoders

## Supervised models

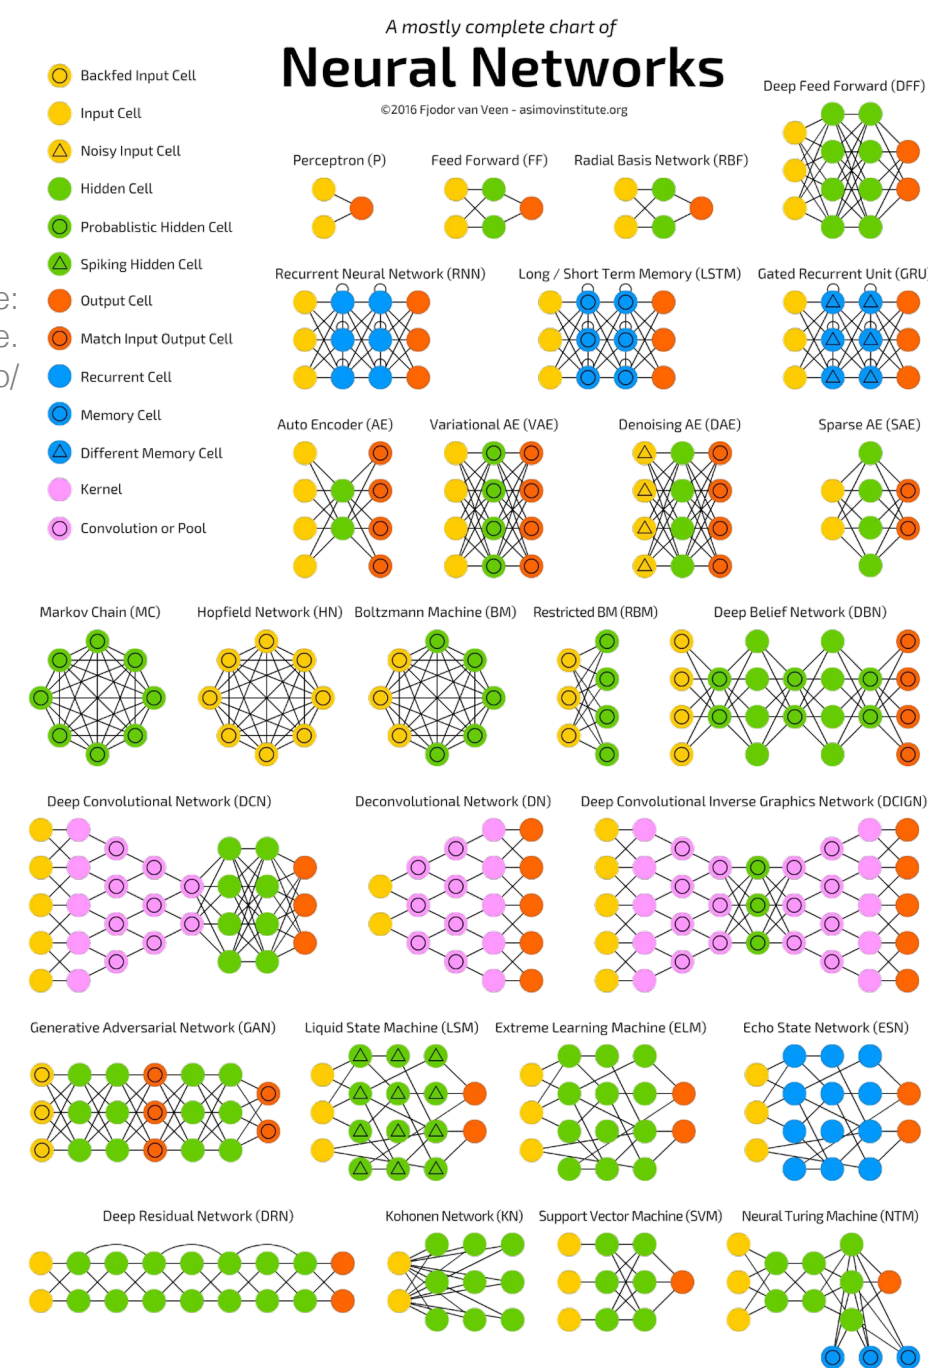
- Image analysis:
  - Convolutional Neural Networks
- Timeseries analysis (often NLP):
  - Recurrent Neural Networks (e.g. LSTMs)
  - Transformer models

## Generative models

- Generative Adversarial Networks (GANs)
- Diffusion Models (e.g. DALL-E 2, Stable Diffusion)
- Generative Pre-trained Transformer (GPT)

# Types of Deep Learning Tools

Azimov Institute:  
<http://www.asimovinstitute.org/neural-network-zoo/>



# Autoencoders


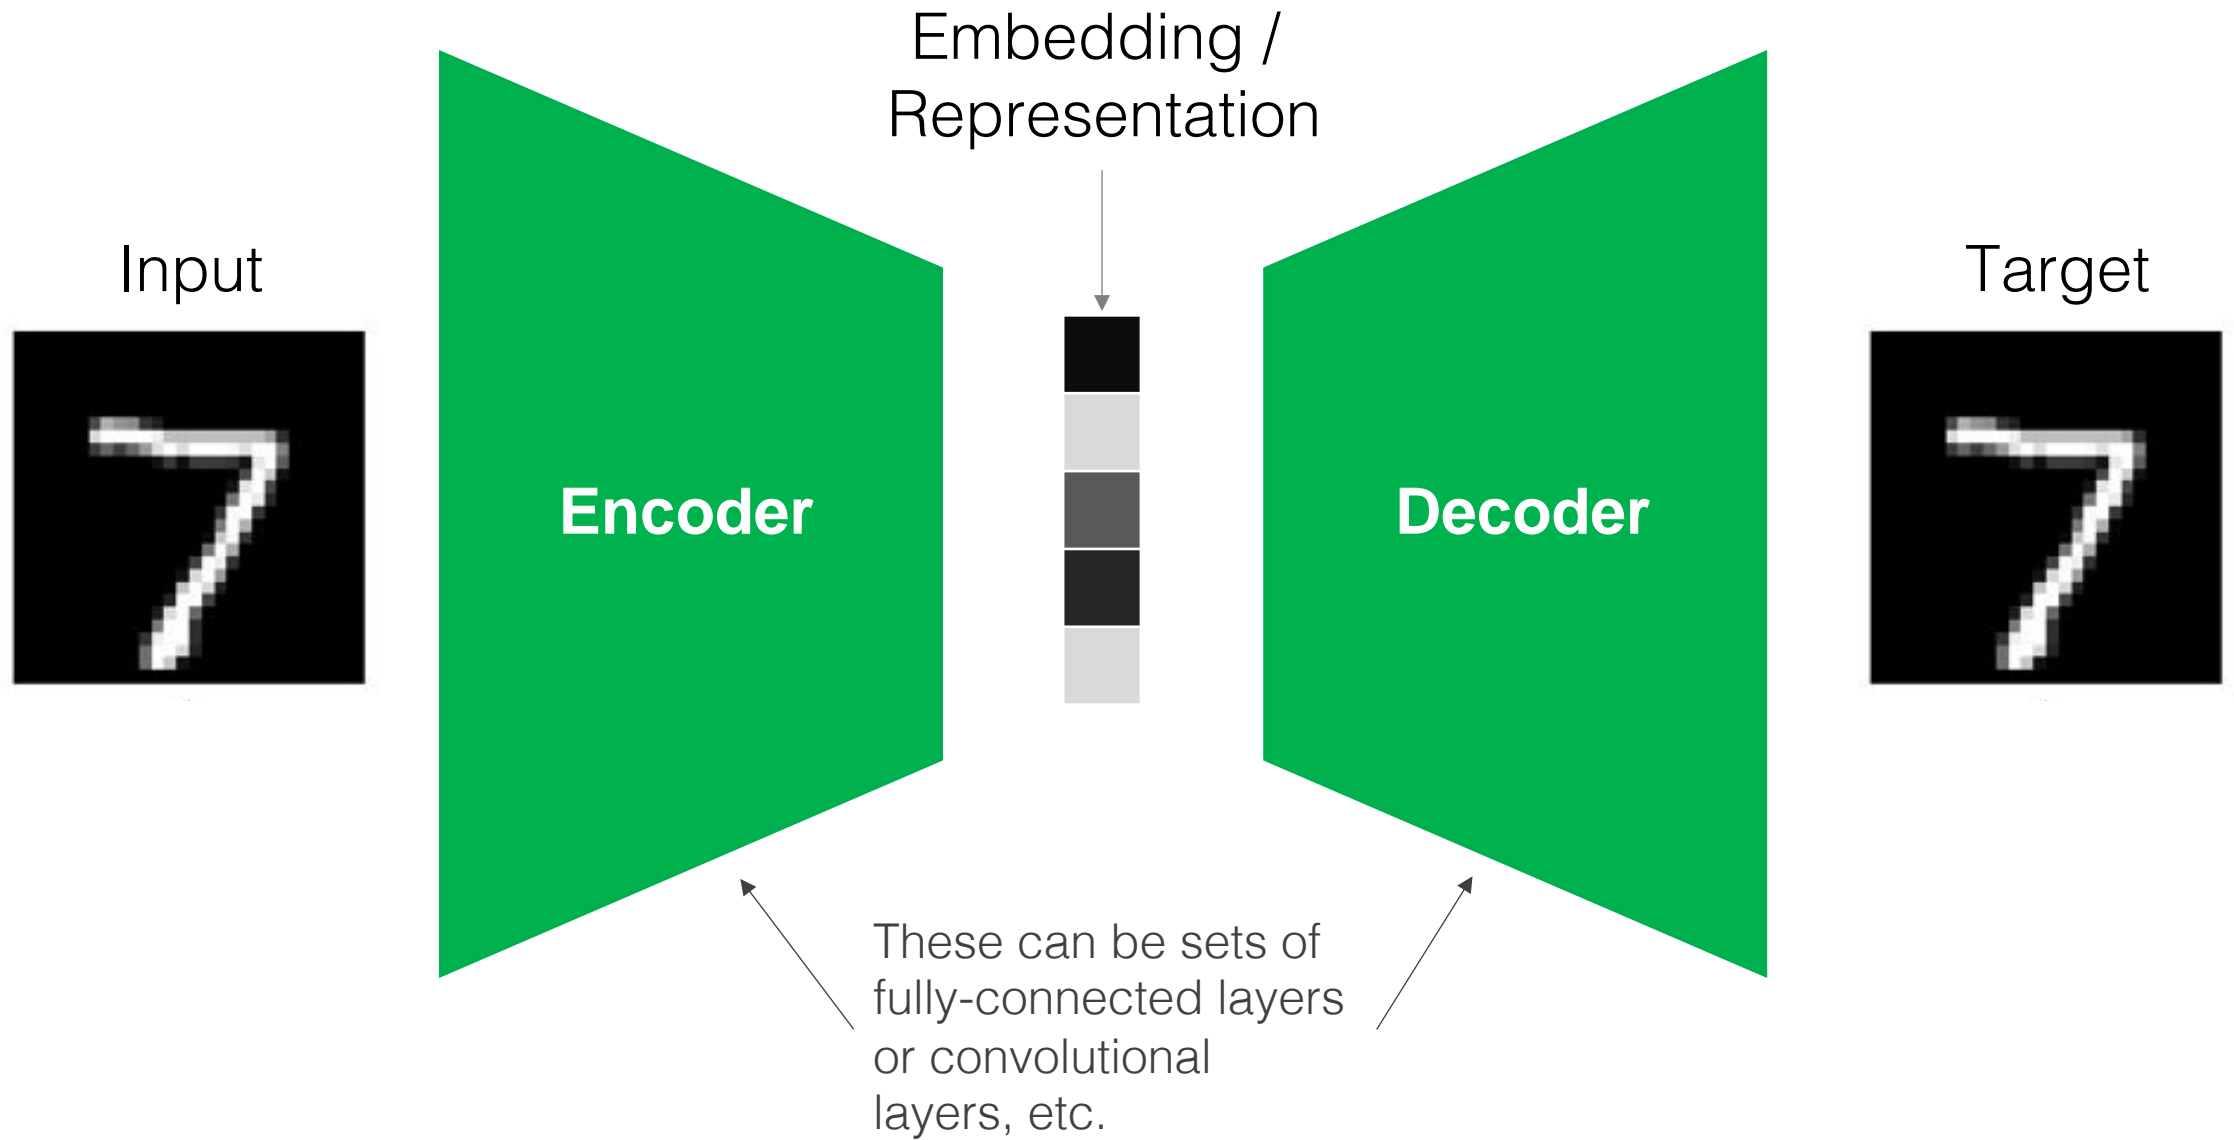


Image from: <https://www.jeremyjordan.me/autoencoders/>



# Autoencoders

Our goal is often to develop a good **encoder** that represents our features well




# Convolutional Neural Networks

# AlexNet



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

## Key





# Convolutional Neural Networks

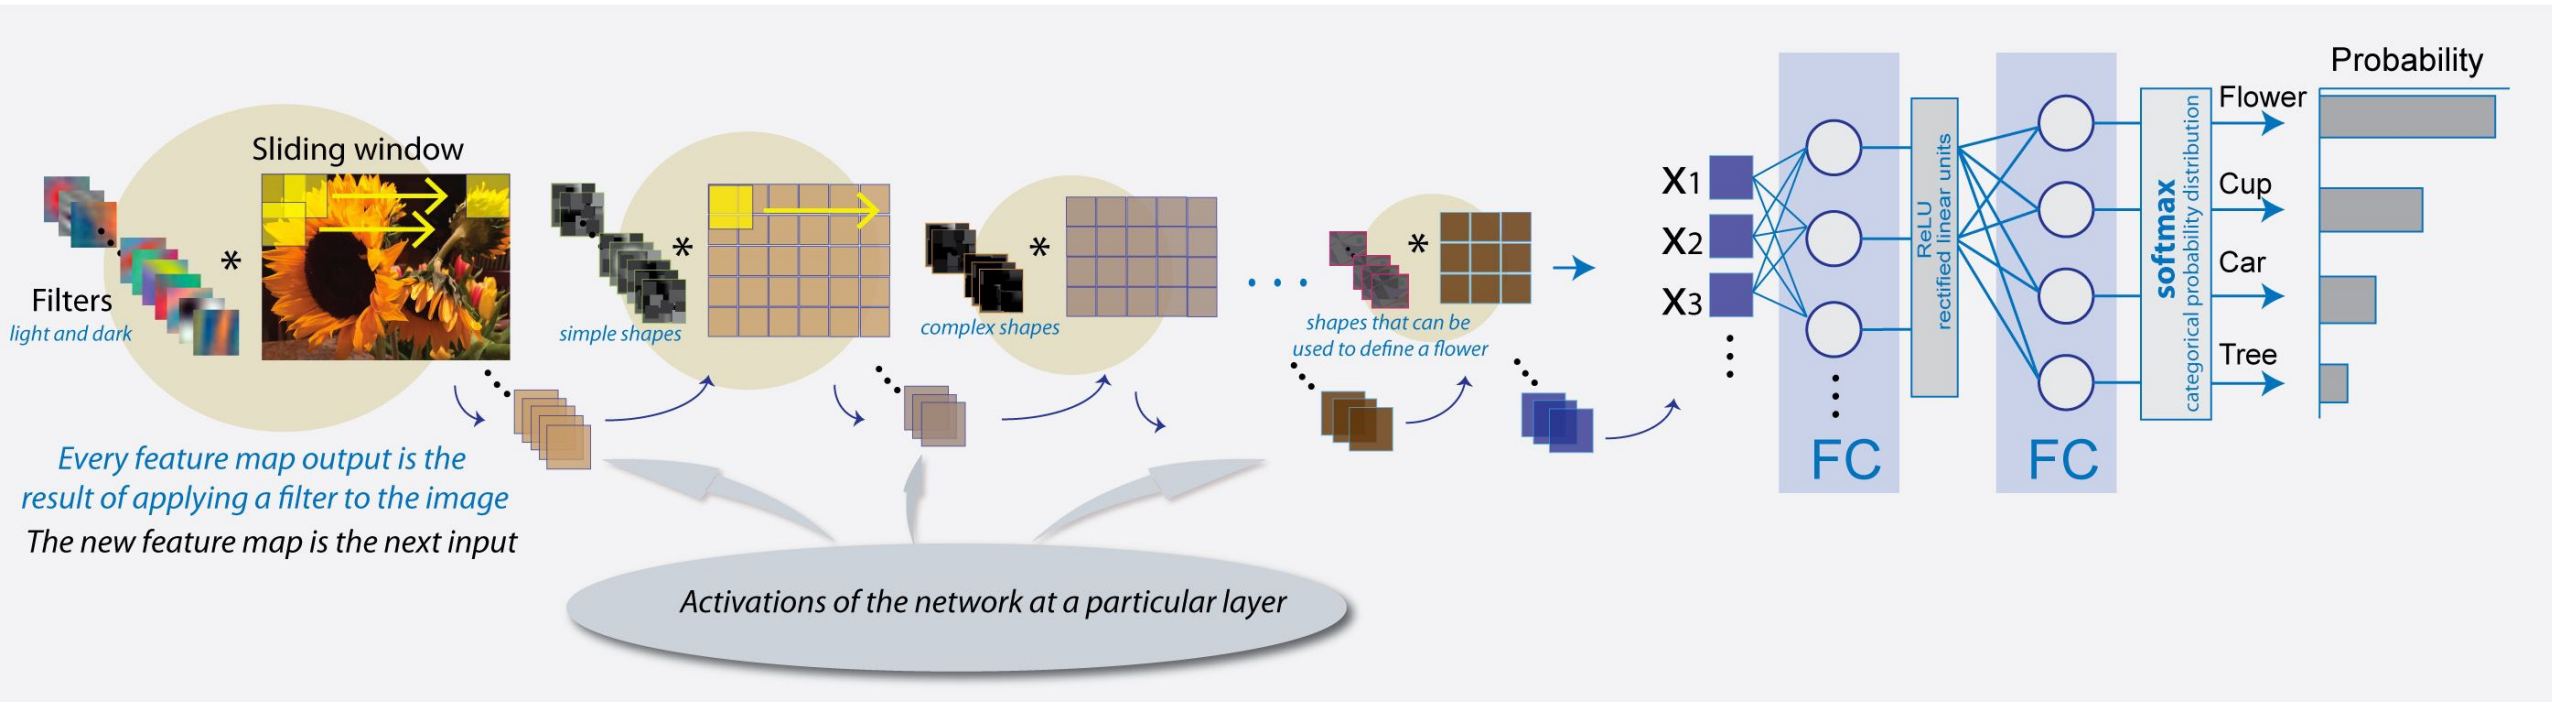
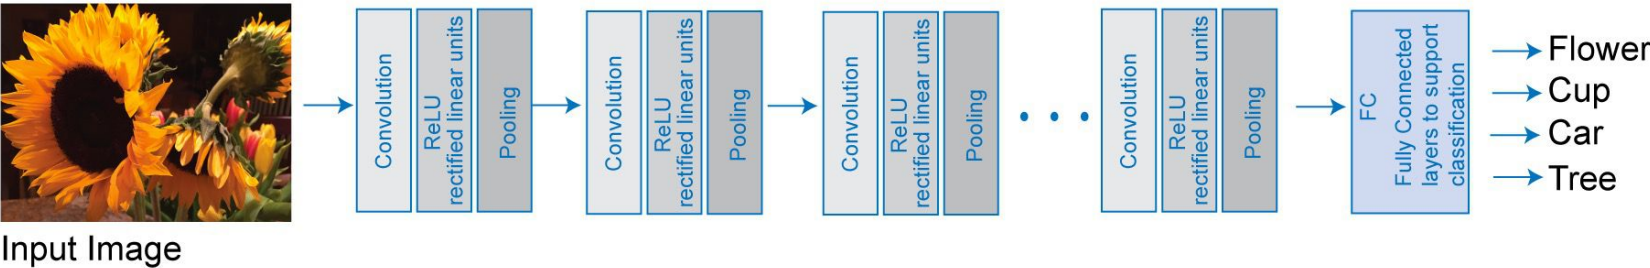


Image from the Mathworks



Data:  $\mathbf{x}$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6



Weights:  $\mathbf{w}$

1	1	1
0	0	0
-1	-1	-1

=

Output:  $\mathbf{x} * \mathbf{w}$


# 2D Convolution

Data:  $x$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $x * w$


Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5
 \end{aligned}$$

# 2D Convolution

Data:  $x$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $x * w$

-6			

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5 = -6
 \end{aligned}$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11		

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 2 + 1 \cdot 5 + 1 \cdot 1 + \\
 &0 \cdot 2 + 0 \cdot 3 + 0 \cdot 2 +
 \end{aligned}$$

$$(-1) \cdot 5 + (-1) \cdot 5 + (-1) \cdot 9 = -11$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

$*$

$=$

Output:  $X * w$

<b>-6</b>	<b>-11</b>	<b>-12</b>	

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 5 & + & 1 \cdot 1 & + & 1 \cdot 4 & + \\
 &0 \cdot 3 & + & 0 \cdot 2 & + & 0 \cdot 0 & +
 \end{aligned}$$

$$(-1) \cdot 5 + (-1) \cdot 9 + (-1) \cdot 8 = -12$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

\*

=

Output:  $X * w$

-6	-11	-12	-11

Computing  
one output  
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 4 + 1 \cdot 2 + \\
 &0 \cdot 2 + 0 \cdot 0 + 0 \cdot 0 +
 \end{aligned}$$

$$(-1) \cdot 9 + (-1) \cdot 8 + (-1) \cdot 1 = -11$$

# 2D Convolution

Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

$*$

$=$

Output:  $X * w$

-6	-11	-12	-11
-7			

Computing one output value:

$$1 \cdot 0 + 1 \cdot 2 + 1 \cdot 3 +$$

$$0 \cdot 4 + 0 \cdot 5 + 0 \cdot 5 +$$

$$(-1) \cdot 6 + (-1) \cdot 3 + (-1) \cdot 4 = -7$$

# 2D Convolution



Data:  $X$

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

6 x 6

Weights:  $w$

1	1	1
0	0	0
-1	-1	-1

3 x 3

\*

=

Output:  $X * w$

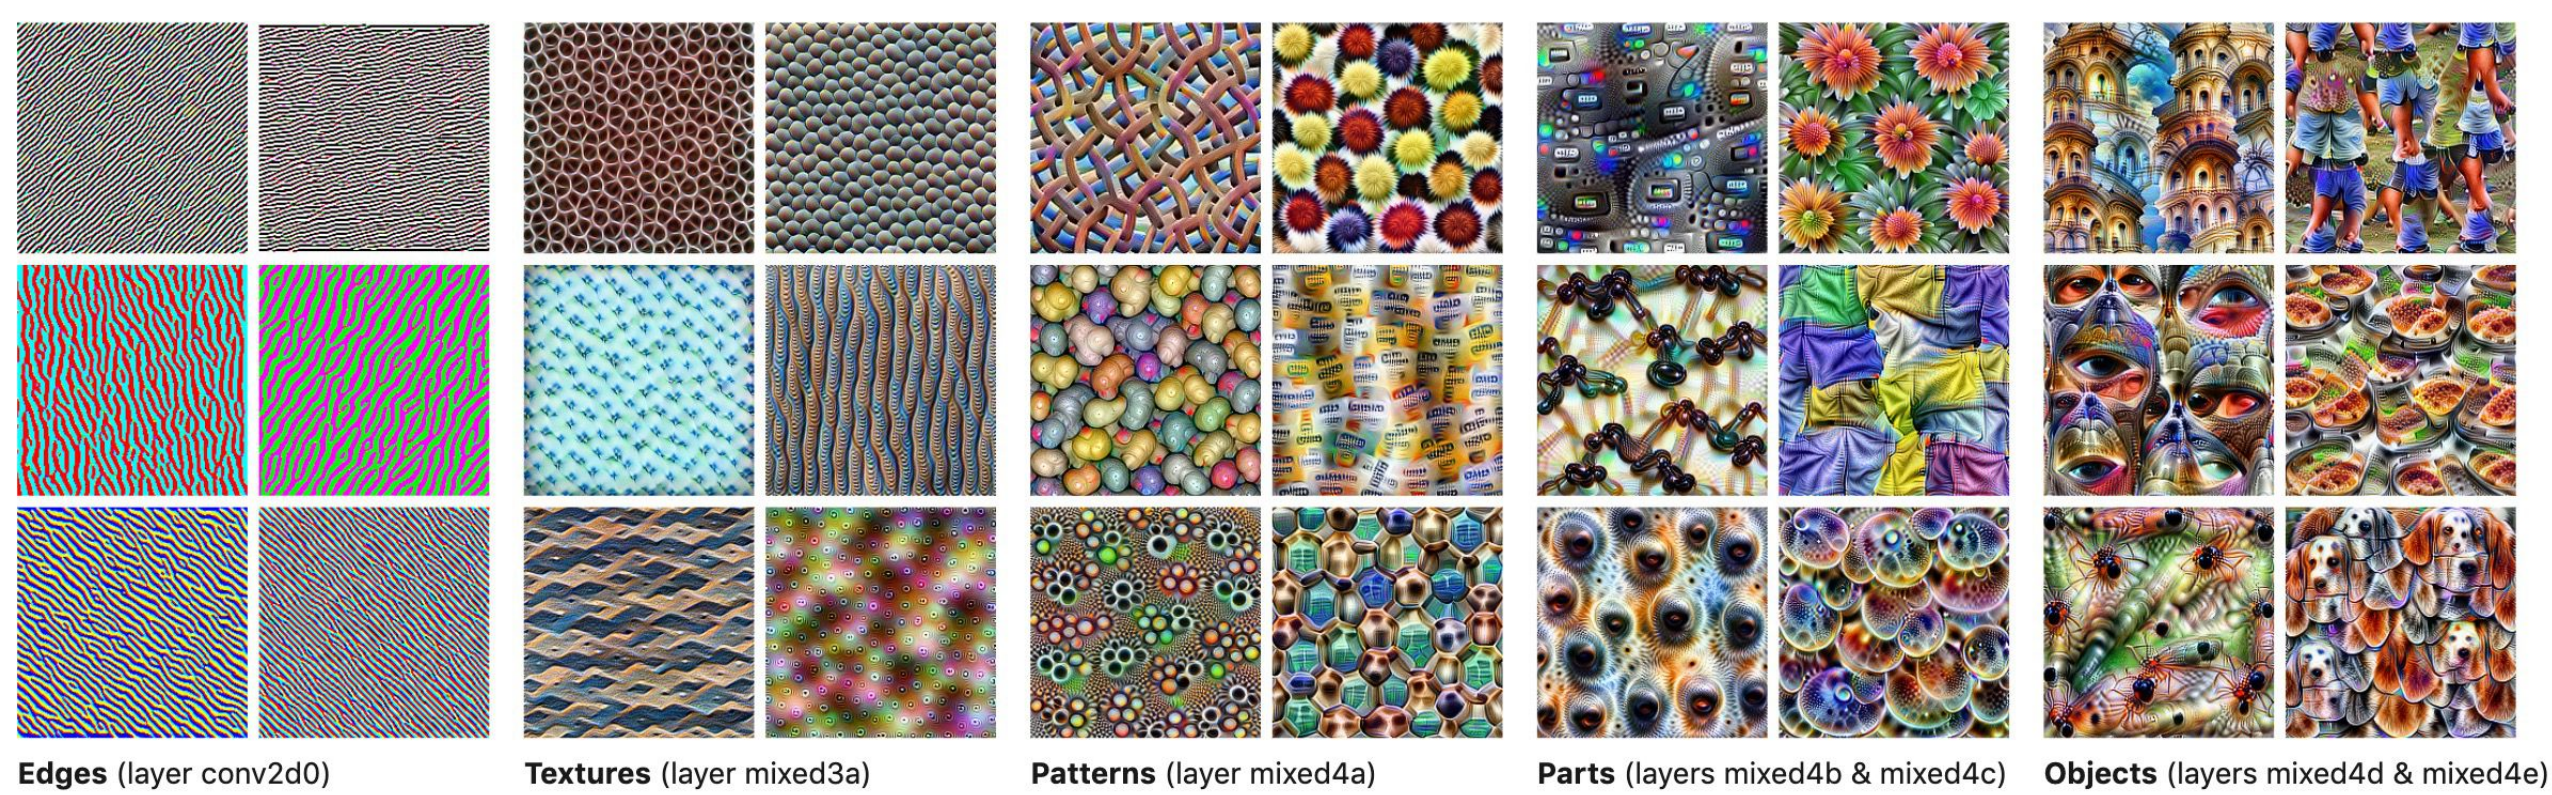

<b>-6</b>	<b>-11</b>	<b>-12</b>	<b>-11</b>
<b>-7</b>	<b>-2</b>	<b>-2</b>	<b>-4</b>
<b>4</b>	<b>1</b>	<b>-2</b>	<b>1</b>
<b>3</b>	<b>-4</b>	<b>-6</b>	<b>-10</b>

4 x 4

# 2D Convolution



# What features do layers respond to?



Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>




# Features


**Dataset Examples** show us what neurons respond to in practice



**Optimization** isolates the causes of behavior from mere correlations. A neuron may not be detecting what you initially thought.



Baseball—or stripes?  
*mixed4a, Unit 6*



Animal faces—or snouts?  
*mixed4a, Unit 240*



Clouds—or fluffiness?  
*mixed4a, Unit 453*



Buildings—or sky?  
*mixed4a, Unit 492*

Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>


# Resources on Visualization of Features

Feature visualization: <https://distill.pub/2017/feature-visualization/>

Building blocks of interpretability: <https://distill.pub/2018/building-blocks/>


Activation Activation Atlases: <https://distill.pub/2019/activation-atlas/>

# Convolution Layer



From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions




$$\text{Parameters} = (5 \times 5 \times 3) \times 6 = 450$$

$$(5 \times 5 \times 6) \times 10 = 1,500$$

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018




# 1 x 1 Convolution Explained




1x1 CONV  
with 32 filters

→  
(each filter has size  
1x1x64, and performs a  
64-dimensional dot  
product)




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

# Max Pooling




max pool with 2x2 filters  
and stride 2




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


# Fully Connected Layer




# AlexNet




Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.




## Key




AlexNet  
(2012)



VGG16  
(2014)



VGG19  
(2014)



Note: an activation function is applied to the output of each layer

Fewer layers,  
larger filters

Key

Input or output layer

Convolutional Layer

Fully Connected Layer


max pooling layer

# CNN Architectures

Adapted from Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018



Scene classification



Object detection




Image segmentation








## Scene classification



AlexNet  
VGG  
GoogLeNet  
ResNet

Inception  
DenseNet  
SqueezeNet  
EfficientNet

Confidence or probability

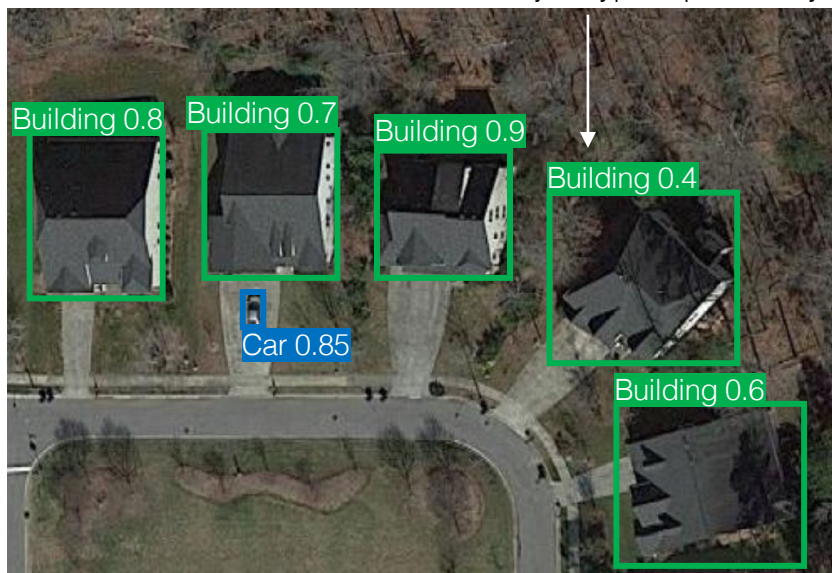


## Object detection



Faster/Fast/R-CNN  
Mask R-CNN  
YOLO  
Single Shot Detector (SSD)  
RetinaNet

Object type & probability



## Image segmentation



U-Net (2015)  
SegNet (2016)  
DeepLab (2017)  
FCN (2016)





# ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

Fei-Fei Li et al. 2010 ([link](#))

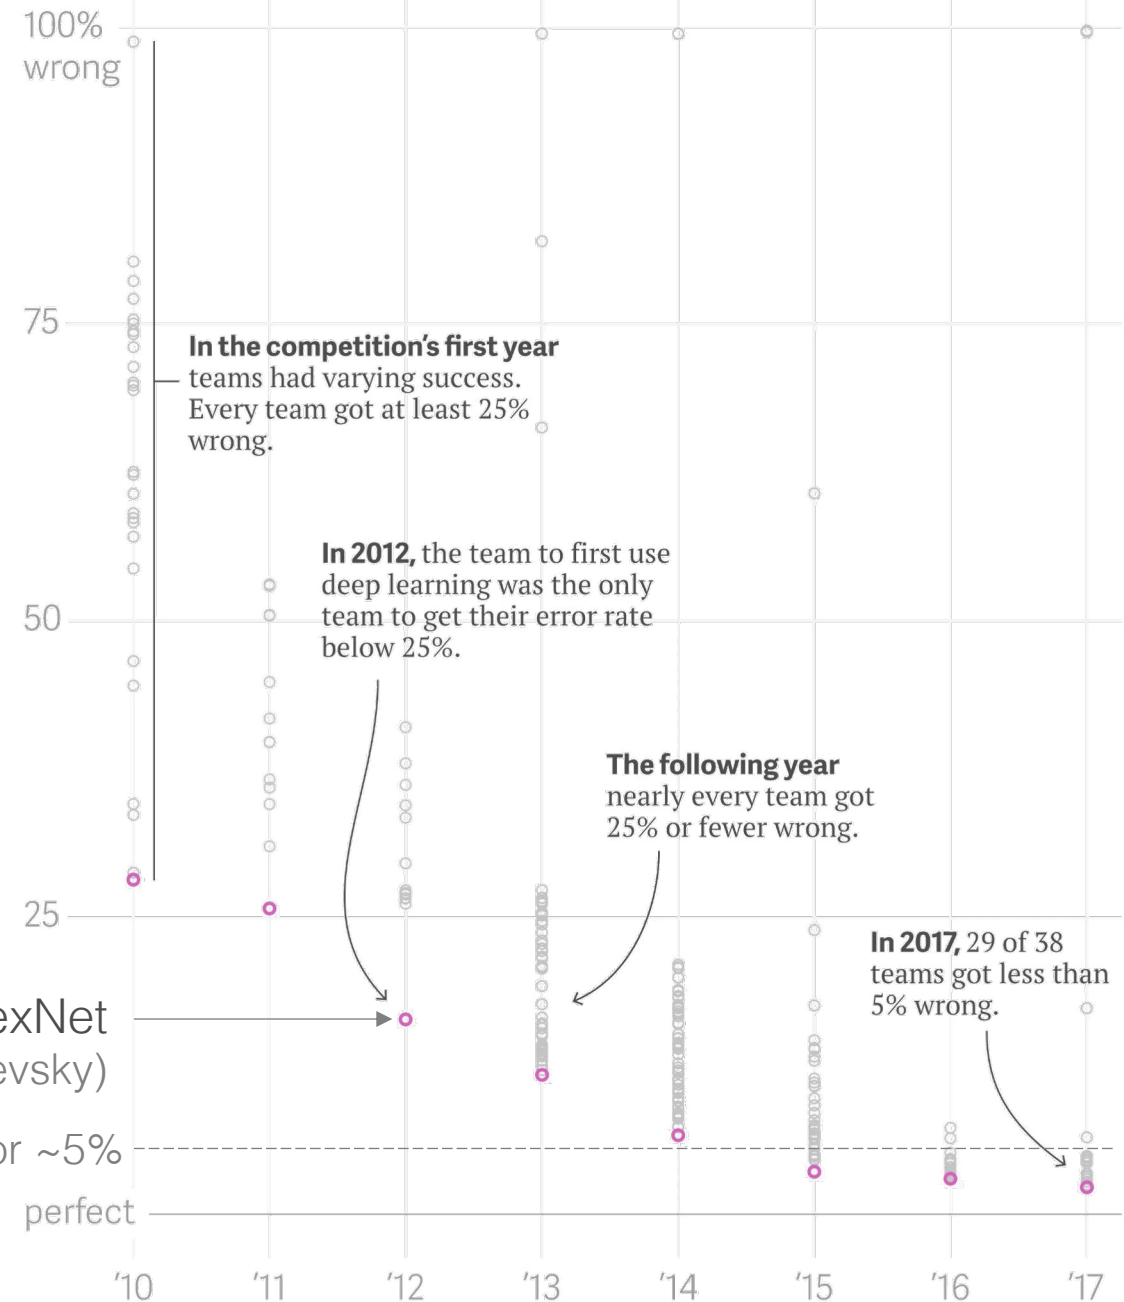
Competition at:  
Conference on Computer Vision and  
Pattern Recognition (CVPR)

## USED FOR MODEL PRETRAINING

AlexNet  
(Hinton, Sutskever, and Krizhevsky)

Human error ~5%



perfect



Source: Quartz, [link](#) David Yanofsky | Quartz

Data: ImageNet

# Deep Learning Models Compared



Models compared for ImageNet

Many of these models are available through Keras ([link](#))

A. Canziani, E. Culurciello and A. Paszke, "Evaluation of neural network architectures for embedded systems," *2017 IEEE International Symposium on Circuits and Systems (ISCAS)*, Baltimore, MD, 2017, pp. 1-4.

# Deep learning frameworks


Tensorflow ([link](#))

Framework for implementing graphical models, such as neural networks



Keras ([link](#))

Wrapper for Tensorflow to make coding easier: higher level and excellent API



Keras

PyTorch ([link](#))

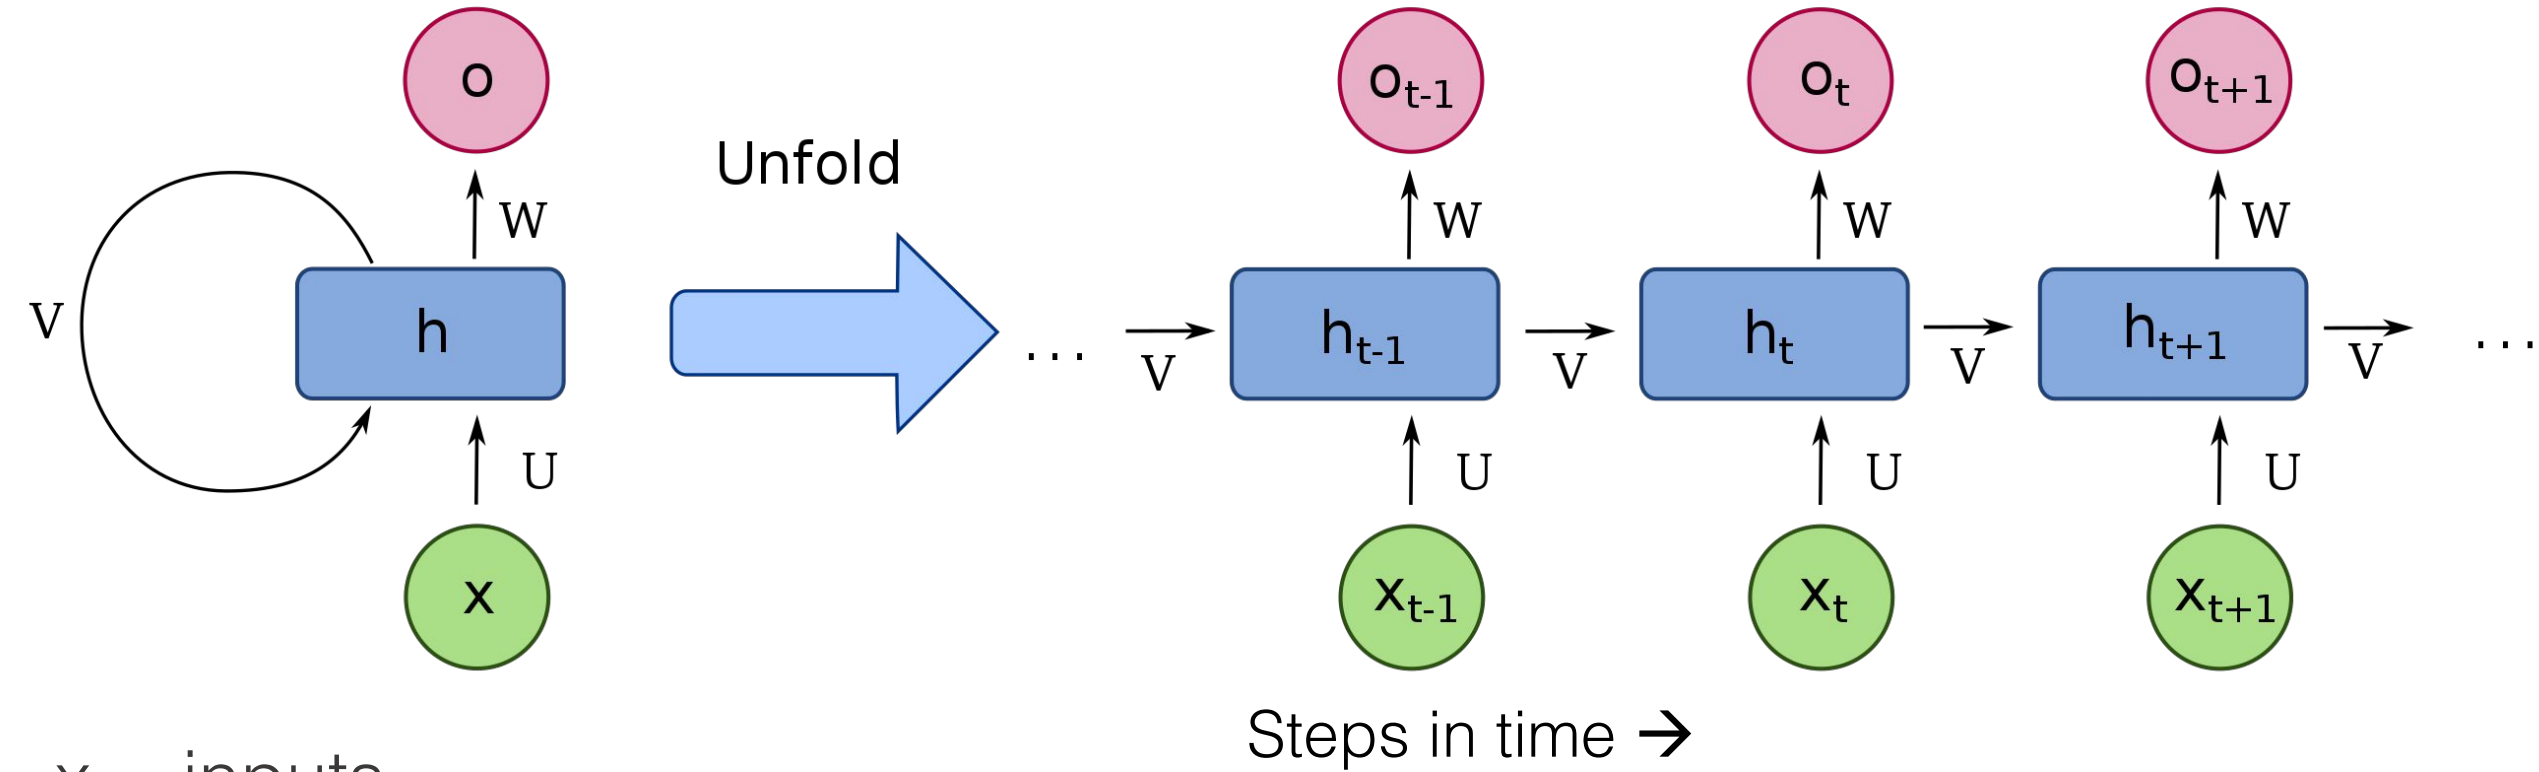
Framework for implementing graphical models, such as neural networks



PyTorch

# KERAS DEMO

# Recurrent Neural Networks



$x$  = inputs

$o$  = outputs

$h$  = hidden layers

$U, V, W$  = model weights

Image from [https://en.wikipedia.org/wiki/Recurrent\\_neural\\_network](https://en.wikipedia.org/wiki/Recurrent_neural_network)

# Generative Adversarial Networks

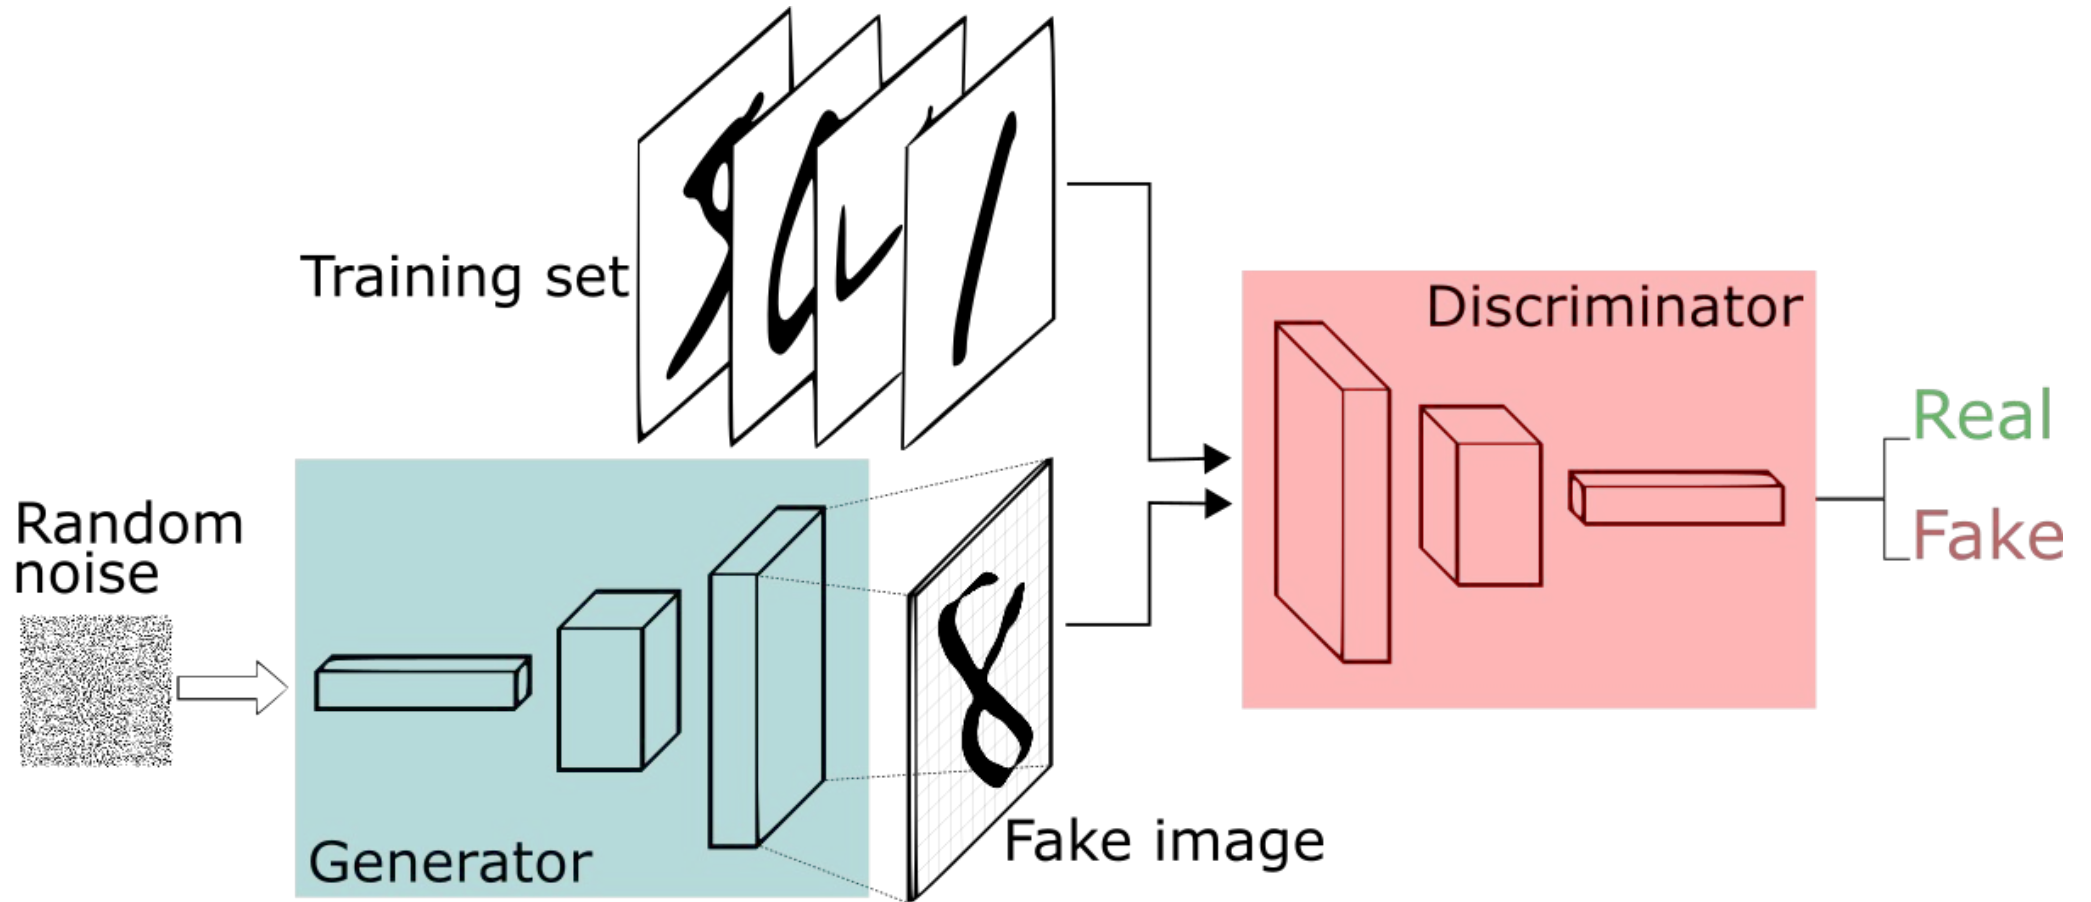


Image from: <https://skymind.ai/wiki/generative-adversarial-network-gan>

# Supervised Learning Techniques

Covered so far

- Linear Regression
- K-Nearest Neighbors
- Perceptron
- Logistic Regression
- Linear Discriminant Analysis
- Quadratic Discriminant Analysis
- Naïve Bayes
- Support Vector Machines
- Decision Trees and Random Forests
- Ensemble methods (bagging, boosting, stacking)
- Neural Networks

Appropriate for:  
● Classification  
● Regression

Can be used with many machine learning techniques