

Reinforcement Learning III

Reinforcement Learning Roadmap

Knowledge of **Environment**

Perfect knowledge

Known Markov
Decision Process



No knowledge

Must learn from
experience

1

Core concepts in reinforcement learning

Actions, Rewards, Value, Environments, and Policies

2

Markov decision processes

...and Markov chains and Markov reward processes

3

Dynamic Programming

How do we find optimal policies?
(Bellman equations)

4

Monte Carlo Control

How do we estimate our value functions?
How do we use the value functions to choose actions?
How do we learn optimal policies from experience?

Markov Decision Process

Components:


State space S

Transition probabilities, P

Rewards, R

Discount rate, γ

Actions, A



Returns (Expected future rewards)

(discount factor weights the the future)

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$$

Policy (how we choose actions)

(can be stochastic or deterministic)

$$\pi(a|s) = P(a|s)$$

State value function

(expected return from state s , and following policy π)

$$v_{\pi}(s) = E[G_t|s]$$

$$v_{\pi}(s) = E[R_{t+1} + \gamma v_{\pi}(S_{t+1})|s]$$

Action value function

(expected return from state s , taking action a , and following policy π)


$$q_{\pi}(s, a) = E[G_t|s, a]$$

$$q_{\pi}(s, a) = E[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1})|s, a]$$

David Silver, UCL, 2015

Bellman Expectation Equations for the **state value** function

(expected return from state s , and following policy π)



$$v_{\pi}(s) = E_{\pi}[G_t | S_t = s]$$

$$\begin{aligned} v_{\pi}(s) &= E_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= E_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \end{aligned}$$

Expectation over the possible actions

Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

$$v_{\pi}(s) = \sum_{\underline{a}} \pi(a|s) \sum_{\underline{s'}} \sum_{\underline{r}} p(s', r | s, a) [r + \gamma v_{\pi}(s')]$$


Bellman Expectation Equations for the **action value** function

(expected return from state s , taking action a , then following policy π)

$$q_{\pi}(s, a) = E_{\pi}[G_t | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) \leftarrow s, a$$



Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

Expectation over the possible actions

$$q_{\pi}(s, a) = \sum_{\underline{s'}} \sum_{\underline{r}} p(s', r | s, a) \left[r + \gamma \sum_{\underline{a'}} \pi(a' | s') q_{\pi}(s', a') \right]$$

Markov Decision Process

If we know the components of the MDP, we can use those to calculate the value functions and determine the optimal policy

Components:

State space \mathcal{S}

Transition probabilities, P

Rewards, R

Discount rate, γ

Actions, \mathcal{A}

Policy (how we choose actions)
(can be stochastic or deterministic)

$$\pi(a|s) = P(a|s)$$

State value function

(expected return from state s , and following policy π)

$$v_{\pi}(s) = E[G_t|s]$$

Action value function

(expected return from state s , taking action a , and following policy π)

$$q_{\pi}(s, a) = E[G_t|s, a]$$

Roadmap to optimal policies

If we assume a **fully known MDP environment**, how do we...
(Markov Decision Process)

1. Evaluate the returns a policy will yield? **Policy evaluation**


2. Find a **better** policy? **Policy improvement**

3. Find the **best** policy? **Policy iteration**

4. Find the best policy **faster**? **Value iteration**


What if we don't have a fully known MDP? **Monte Carlo Methods**

Running example: Gridworld



16 states, 2 of them terminal states labeled “goal”

Valid actions:
(unless there is a wall)



Reward:

-1 for all transitions
(until the terminal state has been reached)

Note: actions that would take the agent off the board are not allowed

Sutton and Barto, 2018

Roadmap to optimal policies

If we assume a **fully known MDP environment**, how do we...
(Markov Decision Process)

1. Evaluate the returns a policy will yield? **Policy evaluation**

2. Find a **better** policy? **Policy improvement**

3. Find the **best** policy? **Policy iteration**

4. Find the best policy **faster**? **Value iteration**

What if we don't have a fully known MDP? **Monte Carlo Methods**

Dynamic Programming

1. Policy Evaluation

Evaluate the returns a policy will yield

Input: policy $\pi(a|s)$

Output: value function $v_\pi(s)$
(unknown)

- 1 Select a policy function to evaluate (estimate the value function)
- 2 Start with a guess of the value function, v_0 (often all zeros)
- 3 **Iteratively** apply the Bellman Expectation Equation to “backup” the values until they converge on the actual value function for the policy, v_π

$$v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_\pi$$

Adapted from David Silver, 2015

1. Policy Evaluation

Evaluate the returns a policy will yield

Policy: $\pi(a|s) = \frac{1}{N_{\text{valid_actions}}}$
for any action a
(i.e. randomly go in any valid direction)

Value function initialization:

$v_0(s) = 0$ for all s (all zeros)

$v_k(s) \rightarrow$ iteration k of policy evaluation

We estimate the value function that corresponds to the policy: $v_\pi(s)$

$v_0(s)$
(initialization)

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

1. Policy Evaluation

Evaluate the returns a policy will yield

Policy: $\pi(a|s) = 1/N_{\text{valid_actions}}$
(randomly go in any direction)

Bellman Expectation Equation:

$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r|s, a) [r + \gamma v_k(s')]$$

1 ($\gamma = 1$)

-1 (rewards are deterministic and constant for all actions)

In Gridworld:

$\frac{1}{N_a}$

1 (once you pick an action
there's no uncertainty as to
which state you'll transition to)

$$v_{k+1}(s) = \sum_a \frac{1}{N_a} (-1 + v_k(s'(a))) = -1 + \sum_a \frac{1}{N_a} v_k(s'(a))$$

Here, the next state is a
deterministic function of a ,
so we can think of it as $s'(a)$

Average of the value of the
 N_a neighboring states

$v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

1. Policy Evaluation

Evaluate the returns a policy will yield

$$v_{k+1}(s) = -1 + \sum_a \frac{1}{N_a} v_k(s'(a))$$

$$v_1 = -1 + \sum_a \frac{1}{4} v_k(s'(a)) = -1$$

$v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

One neighborhood
in $v_0(s)$

	0		
0		0	
	0		

$v_1(s)$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

$v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

 $v_1(s)$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

 $v_2(s)$

0	-1.7	-2	-2
-1.7	-2	-2	-2
-2	-2	-2	-1.7
-2	-2	-1.7	0

 $v_3(s)$

0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0

 $v_{10}(s)$

0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0

 $v_\infty(s) = v_\pi(s)$

0	-14	-20	-22
-14	-18	-20	-20
-20	-20	-18	-14
-22	-20	-14	0

We've found the value function
(expected returns) from our random
movement policy

1. Policy Evaluation

Evaluate the returns a policy will yield

Roadmap to optimal policies

If we assume a **fully known MDP environment**, how do we...
(Markov Decision Process)

1. Evaluate the returns a policy will yield? **Policy evaluation**

2. Find a **better** policy? **Policy improvement**

3. Find the **best** policy? **Policy iteration**

4. Find the best policy **faster**? **Value iteration**

What if we don't have a fully known MDP? **Monte Carlo Methods**

2. Policy Improvement

Find a **better** policy

Input: policy

$$\pi(a|s)$$

Output: better policy

$$\pi'(a|s)$$

Definition of better: has greater or equal expected return in all states:

$$v_{\pi'}(s) \geq v_{\pi}(s) \text{ for all states}$$

- 1 Select a policy function to improve
- 2 Evaluate the value function (our last discussion)
- 3 **Greedy** select a new policy, π' , that chooses actions that maximize value

$$\pi'(s) = \arg \max_a q_{\pi}(s, a)$$

$q_{\pi}(s, a)$ = expected return
from state s , taking action a ,
and following policy π

i.e. pick the **action** that yields the **highest expected returns**

Adapted from David Silver, 2015

Value function:

In this case,
 $q_{\pi}(s, \pi(s)) = v_{\pi}(s)$
 since each action
 leads to only one state

$v_0(s)$


0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$v_{\infty}(s) = v_{\pi}(s)$

0	-14	-20	-22
-14	-18	-20	-20
-20	-20	-18	-14
-22	-20	-14	


Initial policy: $\pi(s)$

$\pi(a|s) =$ randomly go
in any valid
direction



Improved policy
 (in this case this is an
 optimal policy)

$\pi'(s)$



2. Policy Improvement

Find a **better** policy

Roadmap to optimal policies

If we assume a **fully known MDP environment**, how do we...
(Markov Decision Process)

1. Evaluate the returns a policy will yield? **Policy evaluation**

2. Find a **better** policy? **Policy improvement**

3. Find the **best** policy? **Policy iteration**


4. Find the best policy **faster**? **Value iteration**

What if we don't have a fully known MDP? **Monte Carlo Methods**

3. Policy Iteration

Find the **best** policy

Policy **Evaluation**



Policy **Improvement**


⋮

This process will converge to the optimal functions



Input: policy $\pi(a|s)$
Output: **best** policy $\pi^*(a|s)$

Best in the sense that: $v_{\pi^*}(s) \geq v_\pi(s)$ for all states and for all **policies**



Adapted from David Silver, 2015 and Sutton and Barto, 1998

3. Policy Iteration

Find the **best** policy


Input: policy

$\pi(a|s)$

Output: **best** policy

$\pi^*(a|s)$

Policy **Evaluation**




$\text{greedy}(v_\pi) \rightarrow \pi'$
Policy **Improvement**

- 1 Policy Evaluation:** estimate v_π
Iterative policy evaluation
Note: This is VERY slow
- 2 Policy Improvement:** generate $\pi' \geq \pi$
Greedy policy improvement
- 3** Iterate 1 and 2 until convergence

Adapted from David Silver, 2015 and Sutton and Barto, 1998


1 $\pi_0(s)$



$v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

3 $\pi_1(s) = \pi^*(s)$



$v_0(s)$

0	-14	-20	-22
-14	-18	-20	-20
-20	-20	-18	-14
-22	-20	-14	

Policy Evaluation
 $v_\infty(s) \rightarrow v_{\pi_0}(s)$
 Policy Improvement
 2

Policy Evaluation
 $v_\infty(s) \rightarrow v_{\pi_0}(s)$
 4 $v_{\pi^*}(s)$

0	-1	-2	-3
-1	-2	-3	-2
-2	-3	-2	-1
-3	-2	-1	0

Roadmap to optimal policies

If we assume a **fully known MDP environment**, how do we...
(Markov Decision Process)

1. Evaluate the returns a policy will yield? **Policy evaluation**


2. Find a **better** policy? **Policy improvement**

3. Find the **best** policy? **Policy iteration**

4. Find the best policy **faster**? **Value iteration**

What if we don't have a fully known MDP? **Monte Carlo Methods**

$$\pi_0(s)$$



$$v_0(s)$$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$$v_1(s)$$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

What if we stopped after one sweep. This is...

4. Value Iteration

Find the best policy **faster**

4. Value Iteration

Find the best policy **faster**


Input: policy

$\pi(a|s)$

Output: **best** policy

$\pi^*(a|s)$

Policy **Evaluation**



$\text{greedy}(v_\pi) \rightarrow \pi'$
Policy **Improvement**

- 1 Policy Evaluation:** estimate v_π
One-sweep of policy evaluation
- 2 Policy Improvement:** generate $\pi' \geq \pi$
Greedy policy improvement
- 3** Iterate 1 and 2 until convergence

Adapted from David Silver, 2015 and Sutton and Barto, 1998

$v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

 $v_1(s)$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

 $v_2(s)$

0	-1.7	-2	-2
-1.7	-2	-2	-2
-2	-2	-2	-1.7
-2	-2	-1.7	0

 $v_3(s)$

0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0

 $v_{10}(s)$

0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0

 $v_\infty(s) = v_\pi(s)$

0	-14	-20	-22
-14	-18	-20	-20
-20	-20	-18	-14
-22	-20	-14	

So far, we've run policy evaluation all the way to convergence (**this is slow**)

1 $v_0(s)$

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

3 $v_1(s)$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

5 $v_2(s)$

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-1
-2	-2	-1	0

7 $v_3(s) = v_{\pi^*}(s)$

0	-1	-2	-3
-1	-2	-3	-2
-2	-3	-2	-1
-3	-2	-1	0

2 $\pi_0(s)$

	←→	←→	↖
↑	↑↔	↑↔	↑↔
↑	↑↔	↑↔	↑↔
↙	↙↔	↙↔	

4 $\pi_1(s)$

	←	←→	↖
↑	↑↔	↑↔	↑↔
↑	↑↔	↑↔	↓
↙	↙↔	→	

6 $\pi_2(s)$

	←	←	↖
↑	↖	↑↔	↓
↑	↑↔	↘	↓
↙	→	→	

8 $\pi_3(s) = \pi^*(s)$

	←	←	↖
↑	↖	↑↔	↓
↑	↑↔	↘	↓
↙	→	→	


Generalized Policy Iteration

Input: policy

$\pi(a|s)$

Output: **best** policy

$\pi^*(a|s)$



- 1 Policy Evaluation:** estimate v_π
Any policy evaluation algorithm
- 2 Policy Improvement:** generate $\pi' \geq \pi$
Any policy improvement algorithm
- 3** Iterate 1 and 2 until convergence

Adapted from David Silver, 2015 and Sutton and Barto, 1998

Demo

https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_dp.html

So far, we've assumed full knowledge of the environment (MDP)

What if we **DO NOT assume full knowledge of the environment** (MDP)

This means we have to **learn by experience!**

Reinforcement Learning Roadmap

Knowledge of **Environment**

Perfect knowledge

Known Markov
Decision Process



No knowledge

Must learn from
experience

1

Core concepts in reinforcement learning

Actions, Rewards, Value, Environments, and Policies

2

Markov decision processes

...and Markov chains and Markov reward processes

3

Dynamic Programming

How do we find optimal policies?
(Bellman equations)

4

Monte Carlo Control

How do we estimate our value functions?
How do we use the value functions to choose actions?
How do we learn optimal policies from experience?